



# EXPERIENCES WITH MPI RMA AS A FOUNDATIONAL COMMUNICATION ABSTRACTION FOR ONE-SIDED PROGRAMMING MODELS

JEFF HAMMOND, PRINCIPAL ENGINEER


# OUTLINE

- Background on MPI-3 RMA
- Summary of efforts to use MPI-3 RMA
- Overview of ARMCI-MPI
- Overview of OSHMPI v1
- My thoughts on the past, present and future of RMA

## MOTIVATION FOR USING MPI


- Most programmers have better things to do than debug *runtime system issues*. HPC ubiquity requires things to just work.
- *Complex software* needs a language-agnostic\* and language-interoperable programming model / runtime.
- Very few HPC applications bottleneck in communication, so compromising on performance to get portability is a good idea.

\* C/C++, Fortran (3x), Python, C++, Java, C#, D, Go, Perl, Ruby, Rust, Julia, Ocaml, Haskell, Pascal, Ada, ... , but apparently not COBOL.



# HISTORICAL CONTEXT

- Prior to MPI-3, there was a lot of debate about MPI versus PGAS, which was a two-sided versus one-sided debate.
- MPI Forum aspired to make MPI-3 RMA suitable for use in the following:
  - SHMEM
  - Global Arrays (or ARMCI)
  - Fortran coarrays
  - UPC
- MPI RMA working group aspired to make RMA suitable for use on the following:
  - Bad networks
  - Good networks
  - *Imaginary networks*



## STATUS OF RMA USAGE (INCOMPLETE)

Model	Project	Status
<i>Global Arrays</i>	<i>ARMCI-MPI</i>	<i>In production for NWChem</i>
<i>OpenSHMEM</i>	<i>OSHMPI</i>	<i>Useful for research, SMPs</i>
OpenSHMEM	OSHMPI v2	OpenSHMEM 1.4 compliant
Fortran coarrays	OpenCoarrays	GCC 5+
Fortran coarrays	Intel Fortran	Supported in releases since ~2015
Fortran coarrays	CAF 2.0	Published
UPC	GUPC	Evaluated but not using
Grappa	Grappa	Prod. w/ P2P+NBC, no RMA
HPX	HPX	Production w/ P2P, no RMA
Chapel	Chapel	Discussed

# Supporting the Global Arrays PGAS Model Using MPI One-Sided Communication

★ James Dinan, Pavan Balaji,  
Jeff R. Hammond  
*Argonne National Laboratory*  
{*dinan,balaji,jhammond*}@*anl.gov*

Sriram Krishnamoorthy  
*Pacific Northwest National Laboratory*  
*sriram@pnnl.gov*



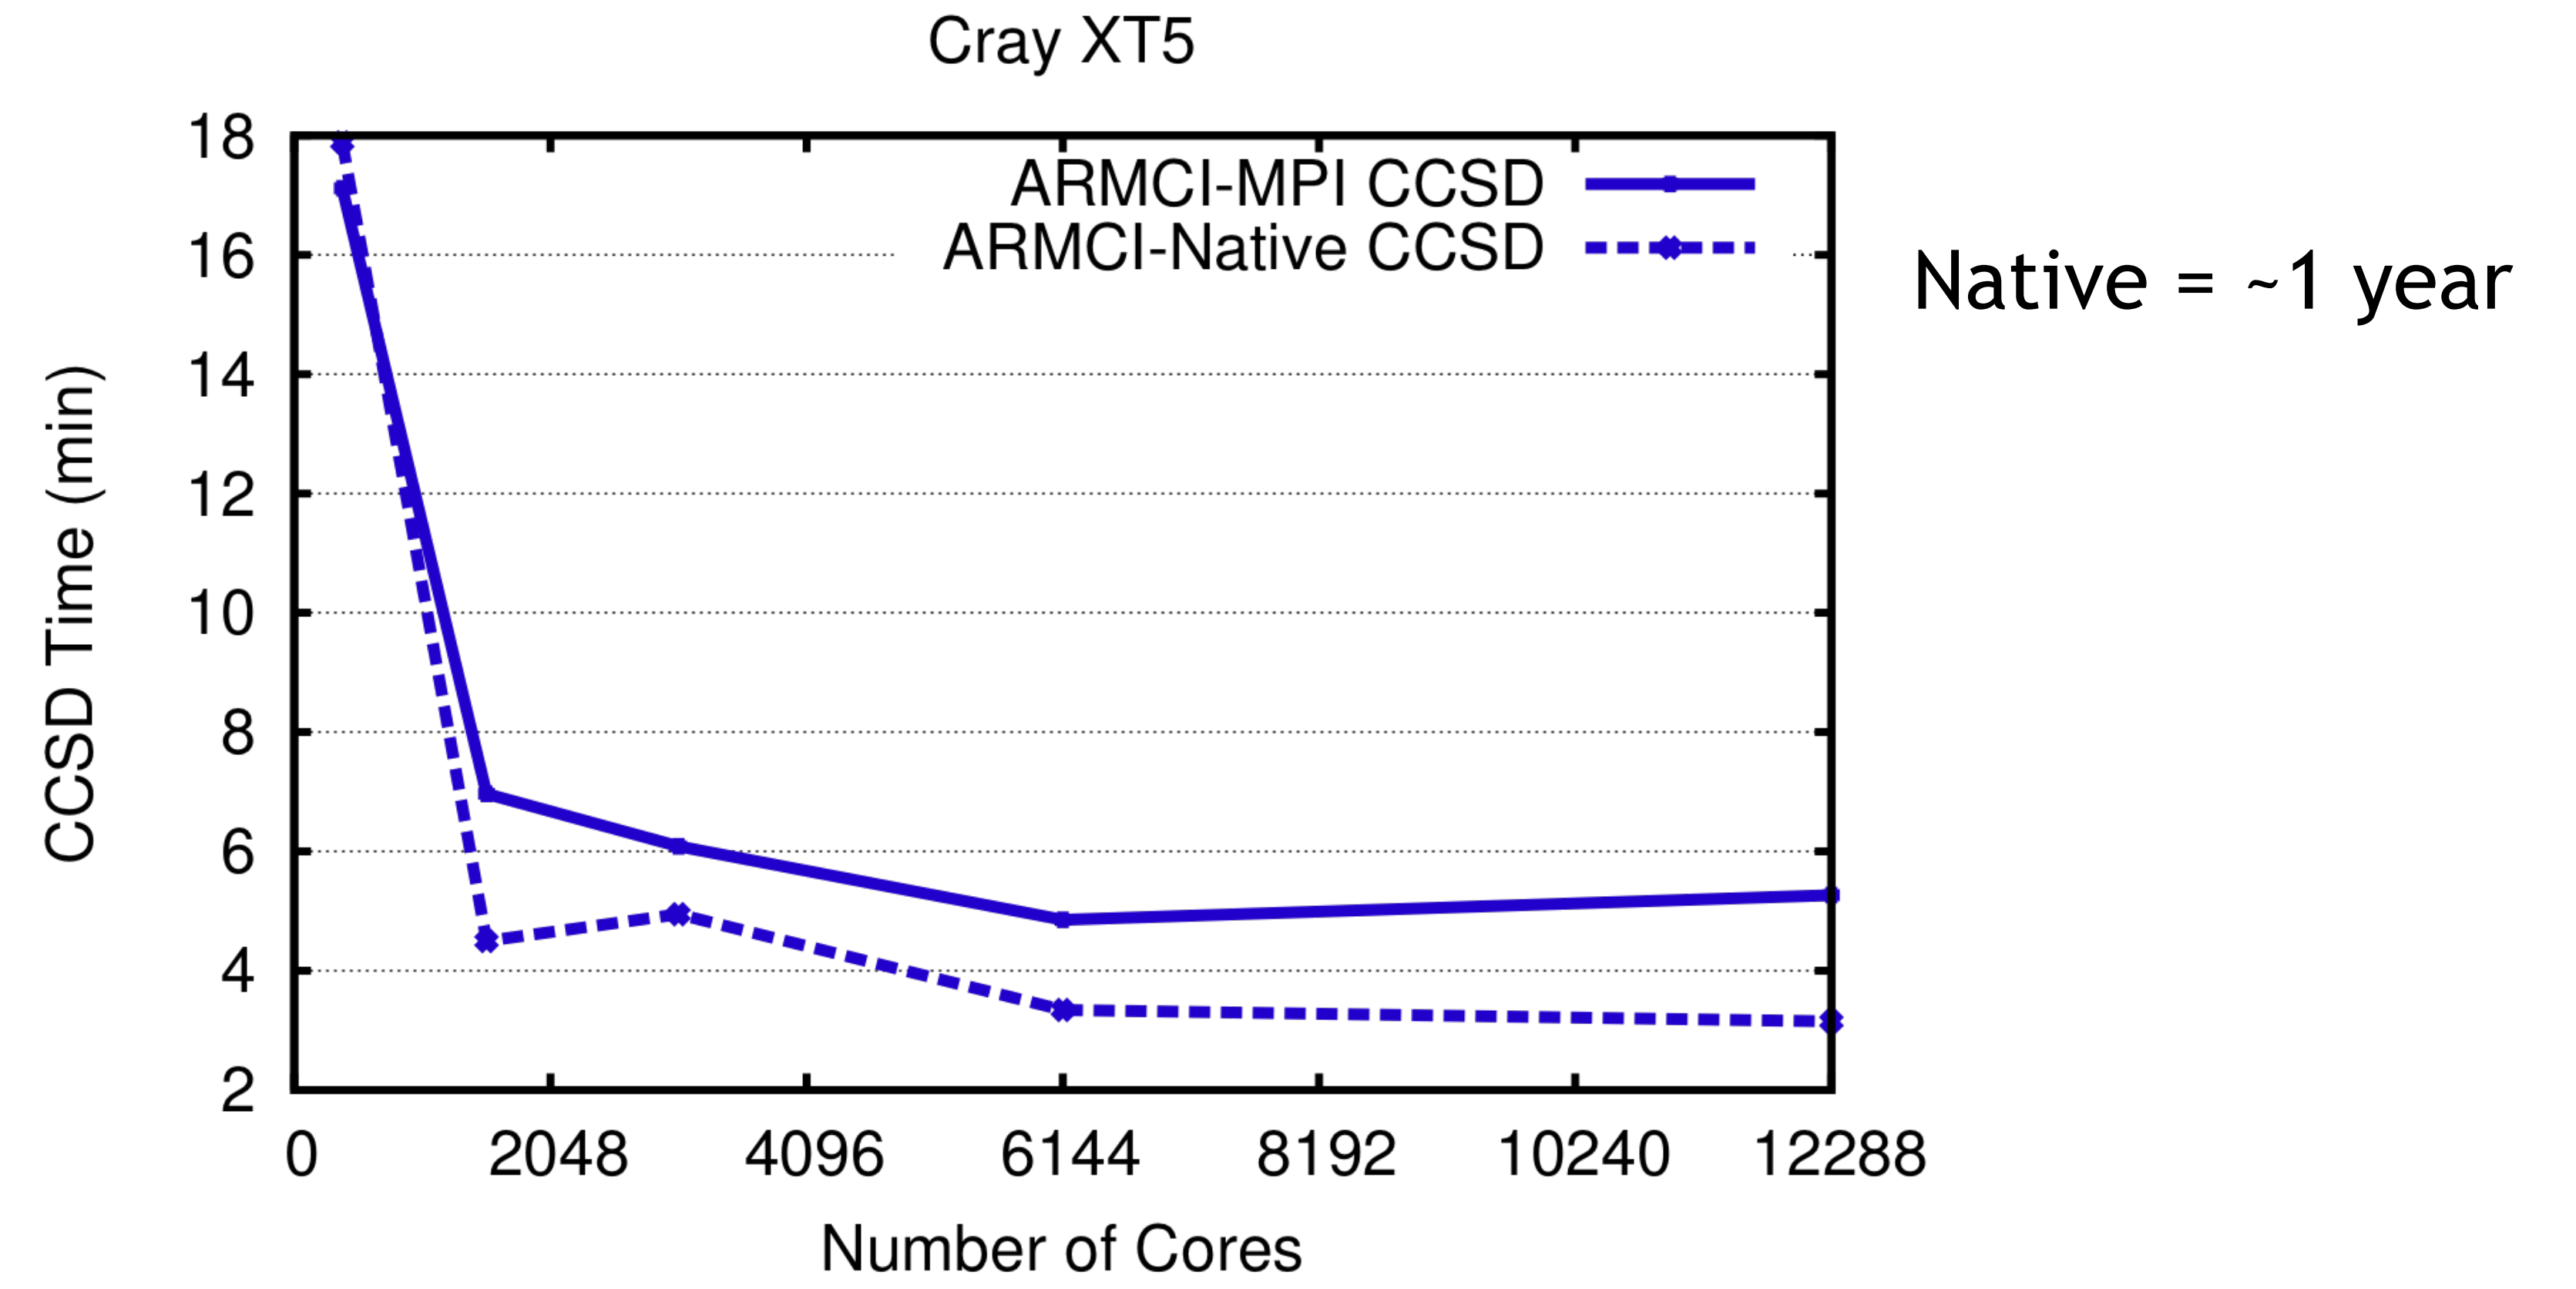

Vinod Tipparaju  
*IEEE Member*<sup>†</sup>  
*tipparajuv@ieee.org*

## ARMCI

- No handles for data, just pointers
- Sequential consistency to same location
- Separate local and remote completion
- Nonblocking RMA
- Atomic operations
- Asynchronous progress guarantee

## MPI-2 RMA

- Opaque handles for data (windows)
- RMA operations unordered
- Local=remote completion
- Nonblocking essentially impossible
- No atomic operations
- No asynchronous progress guarantee



# MPI-3 UPDATE

<https://github.com/pmodels/armci-mpi>

## ARMCI

- No handles for data, just pointers
- Sequential consistency to same location
- Separate local and remote completion
- Nonblocking RMA
- Atomic operations
- Asynchronous progress guarantee

## MPI-3 RMA

- Opaque handles for data (windows)
- *Accumulate operations ordered to same location*
- *Separate local and remote completion*
- *Nonblocking feasible*
- *Atomic operations sufficient*
- No asynchronous progress guarantee



# NWCHEM SCF PERFORMANCE

NWChem 6.3/**ARMCI-MPI3/Casper**

NWChem 6.5/**ARMCI-DMAPP**  
(built by NERSC, Nov. 2014)

iter	energy	time
1	-2830.4366669992	69.6
2	-2831.3734512508	78.8
3	-2831.5712563433	86.9
4	-2831.5727802438	96.1
5	-2831.5727956882	110.0
6	-2831.5727956978	<b>127.8</b>

iter	energy	time
1	-2830.4366670018	67.6
2	-2831.3734512526	85.5
3	-2831.5713109544	105.4
4	-2831.5727856636	126.6
5	-2831.5727956992	161.7
6	-2831.5727956998	<b>190.9</b>

Running on 8 nodes with 24 ppn. Casper uses 2 ppn for comm.

# NWCHEM SCF PERFORMANCE

NWChem 6.3/**ARMCI-MPI3/Casper**

NWChem Dev/**ARMCI-MPIPR**  
(built by NERSC, Sept. 2015)

iter	energy	time
1	-2830.4366669990	69.3
2	-2831.3734512499	77.1
3	-2831.5712604368	84.6
4	-2831.5727804428	93.0
5	-2831.5727956927	107.3
6	-2831.5727956977	<b>128.0</b>

iter	energy	time
1	-2830.4366669999	61.4
2	-2831.3734512509	69.3
3	-2831.5713109521	77.8
4	-2831.5727856618	87.3
5	-2831.5727956974	103.9
6	-2831.5727956980	<b>125.7</b>

Running on 8 nodes with 24 ppn. **Both** use 2 ppn for comm.

# ARMCI-MPI + CASPER

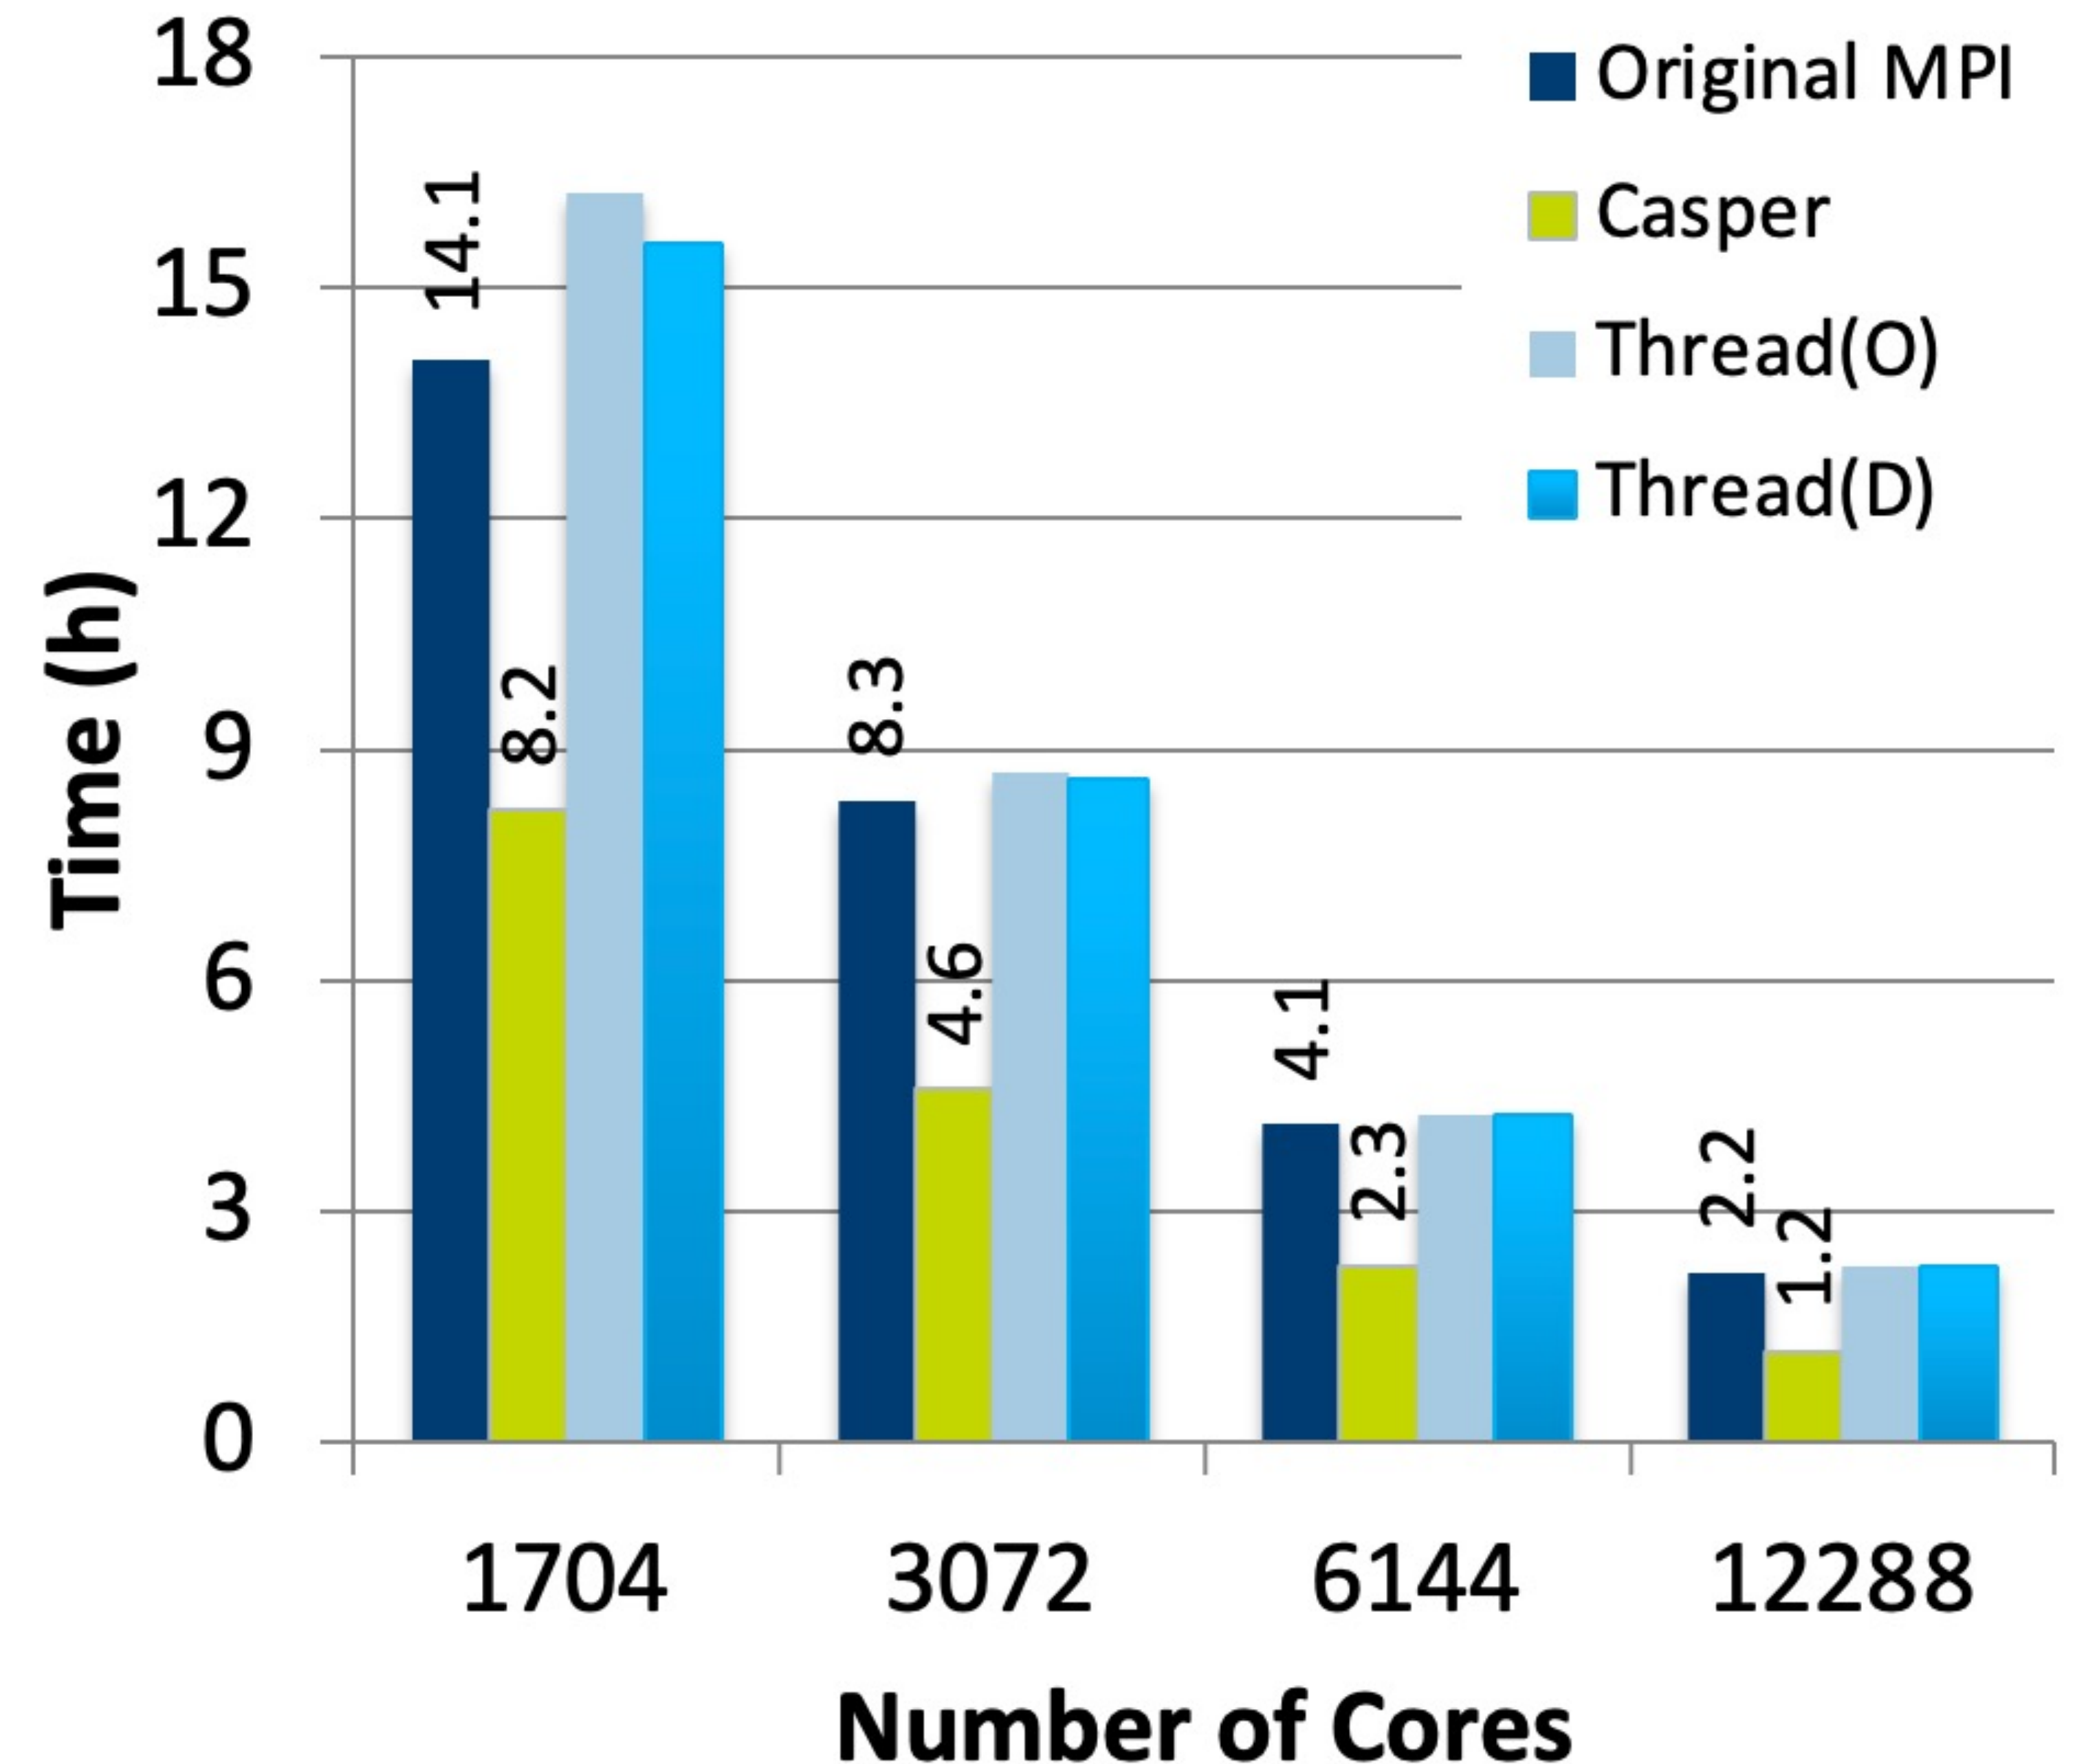
Asynchronous progress is important

## Implementation

- Dedicate process(es) for communication, similar to ARMCI.
- Intercept all RMA calls and redirect using shared-memory (requires Win\_allocate)

## Application usage

- NWChem on 40K+ cores of Cray XC30.
- Bandwidth-limited CCSD(T) for  $(\text{H}_2\text{O})_{21}$ .



# ARMCI-MPI SUMMARY

Made NWChem universally portable:

- Only ran on Blue Gene/Q because of ARMCI-MPI (but MPI-2 RMA ☹️)
- Ran on ARM32 out-of-the-box in 2013
- No weird failures on IB or problems with >>2GB arrays

But

- Asynchronous progress is still a problem. Casper is not always better.
- Latency is too high. ARMCI over two-sided often wins for DFT code.
- Open-MPI correctness issues in RMA are still causing problems for users and developers.

A direct port of Global Arrays would have been better (but much more work)...

# Implementing OpenSHMEM using MPI-3 one-sided communication

Jeff R. Hammond<sup>1</sup>, Sayan Ghosh<sup>2</sup>, and Barbara M. Chapman<sup>2</sup>

## ARMCI-MPI

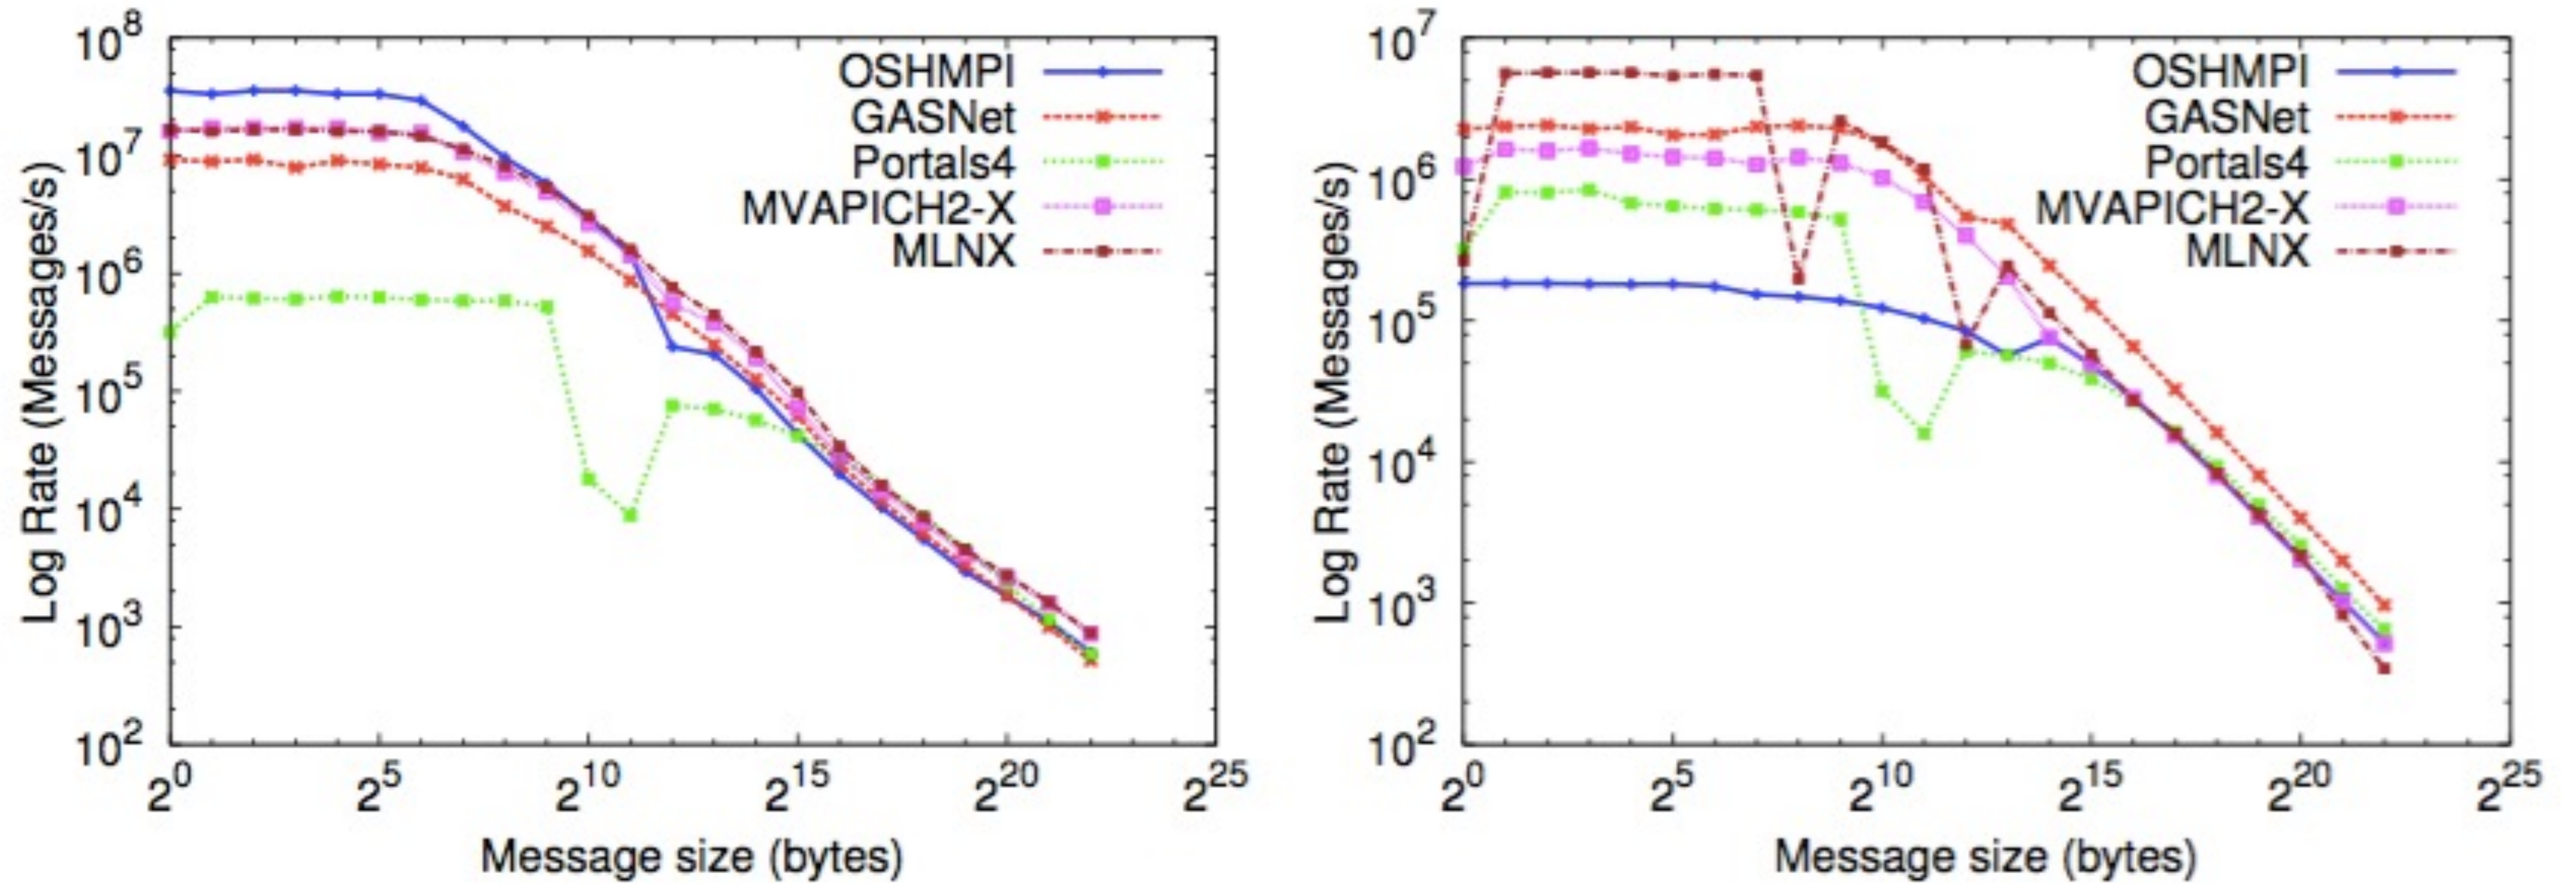
- One window for every allocation (expensive window lookup)
- Reverse-engineered semantics from PNNL implementation
- Weird contortions for ARMCI groups
- Workarounds for asynchronous progress problems

## OSHMPI v1

- Symmetric heap suballocated from one window (fast lookup)
- Direct translation from OpenSHMEM specification
- Weird contortions for SHMEM PE subsets
- Fast path for intranode communication
- Ignores SAME\_OP\_NO\_OP nonsense


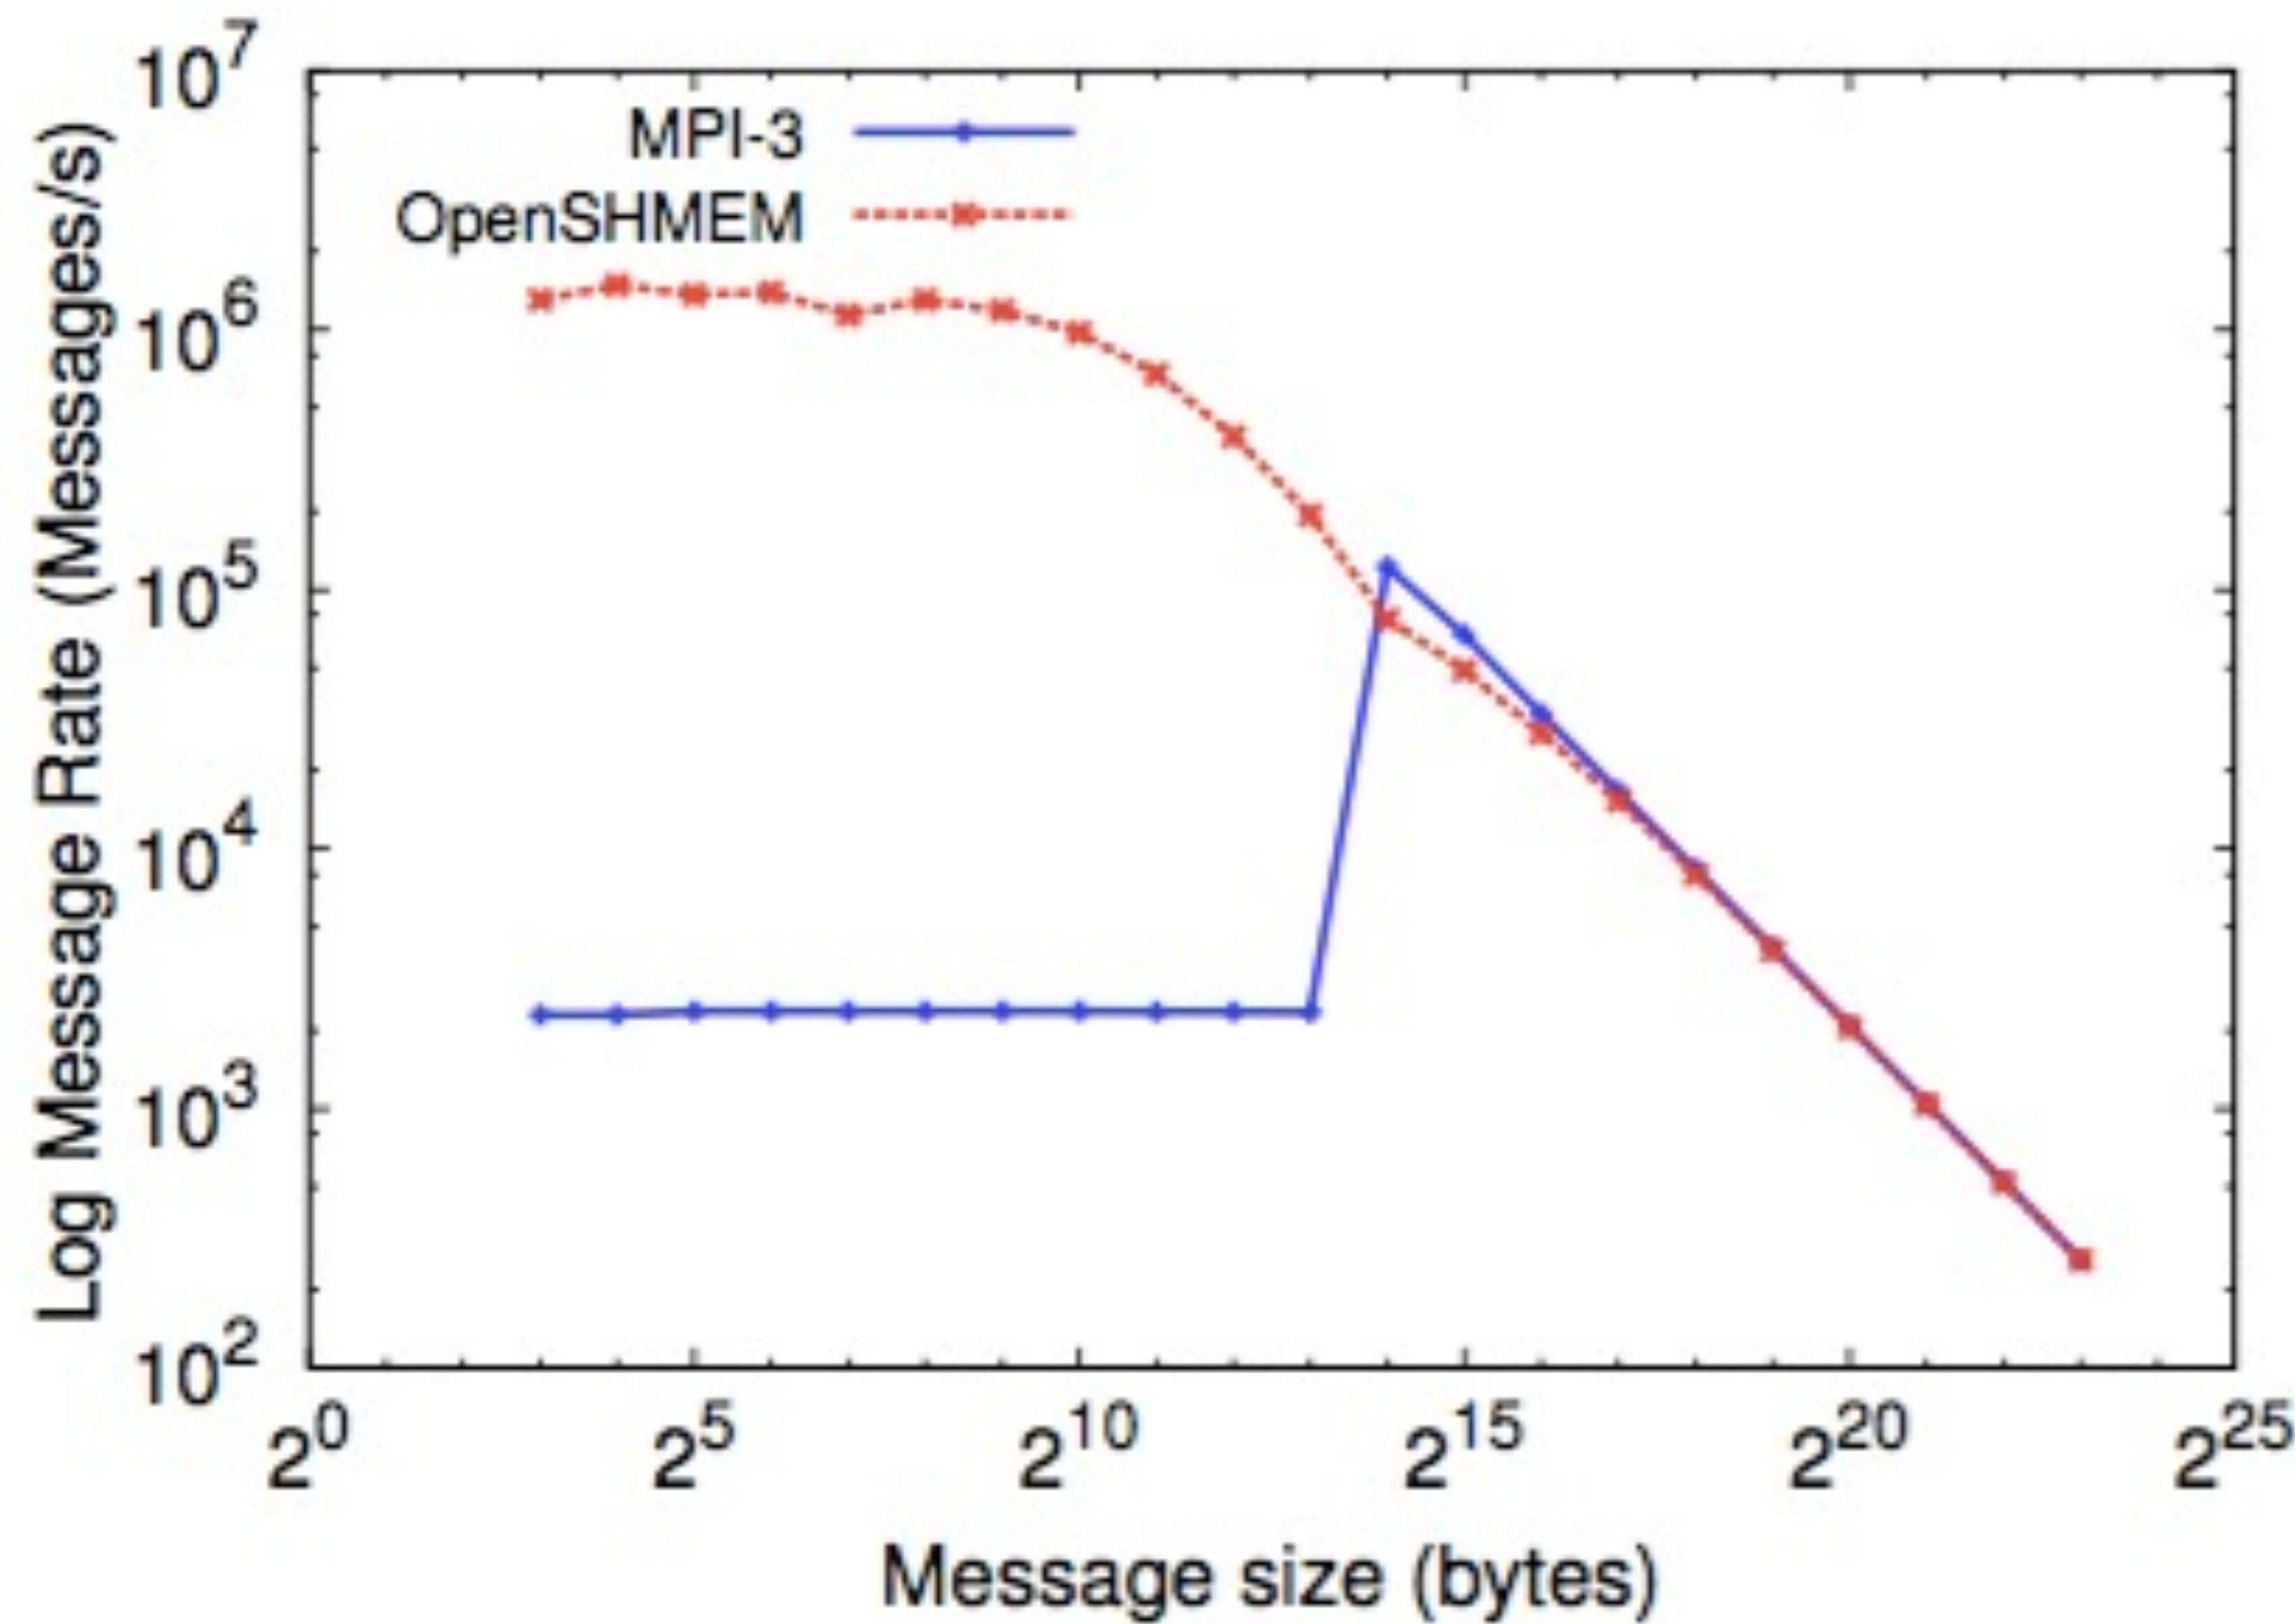
<https://github.com/jeffhammond/oshmpi>

# OSHMPI PUT MESSAGE-RATE ON IB



**Figure:** Intranode (left) and internode (right).

# MPI IMPLEMENTATION OVERHEAD IS (WAS?) HIGH



Internode (left), intranode (right). MVAPICH2-X from 2013-2014.

# OSHMPI SUMMARY

- **Proved that MPI-3 RMA is a viable back-end for OpenSHMEM (hence v2)**
- Easy to install on every platform
- Very good performance in shared memory only because it bypassed RMA altogether
- Very bad performance in distributed memory because of MPI RMA implementations
- Best SHMEM implementation for benchmarks dominated by collectives 😊





**FINAL  
THOUGHT**

# 1. IMPLEMENTATIONS ARE THE PROBLEM

- Implementations continue to be bad at performance:
  - Latency is not good
  - Message-rate is not good
  - Bandwidth is inconsistent
  - Asynchronous progress is almost non-existent
- Implementations continue to be bad at correctness:
  - MPICH (and derivatives) are correct almost all of the time...
  - NWChem = MPI\_Accumulate + MPI\_TYPE\_SUBARRAY + MPI\_SUM + MPI\_DOUBLE
  - Insufficient to test correctness in shared-memory

## 2. RMA IS COMPLICATED AND HARD TO USE




- MPI-1 features were easy to use so scientists assume MPI is easy to use, which is false for RMA
- RMA offers many ways to do the same thing:
  - Windows: allocate, create, dynamic, shared
  - Sync: fence, PSCW, lock, lock\_all, flush, flush\_all, flush\_local, etc.
- The standard should explicitly recommend allocate + lock\_all + flush(\_local) as the preferred RMA motif
- We need a user guide for RMA somewhere, and a set of benchmarks to determine all the platform-specific dependence of RMA features

### 3. RMA NEEDS MINOR FIXES

- Users should be allowed to query shared memory in an allocated window:  
<https://github.com/mpi-forum/mpi-issues/issues/23>
- There should be a request-based version of everything:  
<https://github.com/mpi-forum/mpi-issues/issues/128>
- MPI\_PROD should be removed (there are **zero** use cases and **zero** users)
- The default should be ANY\_OP, not SAME\_OP\_NO\_OP
- Nonblocking remote flush is a good idea

## 4. RMA IS THE BEST MODEL FOR GPU COMMUNICATION

- GPUs are really good at moving data of all sizes
- Synchronization is expensive on GPUs due to massive parallelism
- RMA properly separates data movement from synchronization and supports memory allocation and registration



<https://www.nvidia.com/en-us/data-center/nvlink/>